

IMPORT LIBRARIES AND DATASET

```
In [6]: 1 import plotly.express as px
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import numpy as np
5 import seaborn as sns
6 import matplotlib.pyplot as plt
7 from math import pi
8 import string
9 import re
10 import tensorflow as tf
11 from surprise import Dataset, Reader
12 from sklearn.model_selection import train_test_split
13 from surprise import KNNBasic
14 from surprise.accuracy import rmse
15 from sklearn.preprocessing import StandardScaler
16 from sklearn.neighbors import KNeighborsClassifier
17 from sklearn.ensemble import RandomForestClassifier
18 from sklearn.metrics import roc_curve, auc
19 import matplotlib.pyplot as plt
20 from surprise import Dataset, Reader, KNNBasic
21 from surprise.model_selection import train_test_split
22 from surprise.accuracy import rmse
23 import researchpy as rp
24 from pgmpy.models import BayesianModel
25 from pgmpy.estimators import MaximumLikelihoodEstimator
26 from pgmpy.inference import VariableElimination
27 from nltk.corpus import stopwords
28 from sklearn.linear_model import LogisticRegression
29 from sklearn.neighbors import KNeighborsClassifier
30 from sklearn.model_selection import train_test_split
31 from sklearn.metrics import roc_curve, roc_auc_score
32 from sklearn.preprocessing import LabelEncoder
33 from scipy.stats import chi2_contingency
```

```
34 from matplotlib.colors import Normalize
35 from sklearn.model_selection import train_test_split
36 from sklearn.neighbors import KNeighborsClassifier
37 from sklearn.metrics import accuracy_score
38 from sklearn.preprocessing import MinMaxScaler
39 from sklearn.impute import SimpleImputer
40 import warnings
41 import plotly.graph_objects as go
42 from sklearn.preprocessing import MaxAbsScaler
43 from sklearn.preprocessing import MaxAbsScaler
44 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
45 from sklearn.preprocessing import StandardScaler, LabelEncoder
46 from sklearn.model_selection import GridSearchCV
47 from sklearn.ensemble import RandomForestClassifier
48 from sklearn.preprocessing import LabelEncoder, StandardScaler
49 from sklearn.model_selection import train_test_split, GridSearchCV
50 from sklearn.feature_selection import SelectKBest, f_classif
51 from sklearn.metrics import confusion_matrix, accuracy_score
52 from sklearn.model_selection import cross_val_score
53 import statsmodels.api as sm
54 import squarify
55 import math
56 from sklearn.metrics.pairwise import cosine_similarity
57 from sklearn.impute import SimpleImputer, MissingIndicator
58 from sklearn.model_selection import train_test_split
59 from sklearn.ensemble import RandomForestClassifier
60 from sklearn.metrics import accuracy_score
61 import pandas as pd
62 import numpy as np
63 import pandas as pd
64 import numpy as np
65 from sklearn.model_selection import train_test_split
66 import matplotlib.pyplot as plt
67 import seaborn as sns
68 from sklearn.model_selection import train_test_split
69 from sklearn.neighbors import KNeighborsClassifier
70 from sklearn.metrics import classification_report
71 from sklearn.linear_model import LogisticRegression
72 from sklearn.tree import DecisionTreeClassifier
```

```

74 from sklearn.feature_extraction.text import TfidfVectorizer
75 from sklearn.metrics.pairwise import cosine_similarity
76 from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score
77 # Read the Excel file into a DataFrame
78 data= pd.read_excel(r"C:\Users\Administrator\Desktop\DATA.xlsx")
79 # display the dataframe
80 data.head()

```

Out[6]:

	user ID	Name	Age	Gender	Academic Background	Field of Study	Skills	Industry_Interest	Job_Type_Interest	Location_Interest	Salary_Expectation_(in_USD)	d
0	986206	John Smith	28	Male	Bachelor's degree	Computer Science 2	Java, Python, Data Structures	Technology	Full-time	New York	65000	
1	769632	Jane Doe	42	Female	Master's degree	Business	Finance, Accounting, Microsoft Excel	Finance	Contract	London	85000	
2	981314	David Lee	35	Male	bachelor's degree	NaN	Sales, Customer Service, Communication	Retail	Part-time	Chicago	30000	
3	962892	Sarah Johnson	27	Female	Associate's degree	Nursing	Patient Care, Medical Terminology	Healthcare	Full-time	Los Angeles	45000	
4	967782	Michael Williams	46	Male	Bachelor's degree	Marketing	Digital Marketing, Social Media	Marketing	Freelance	Toronto	70000	

In [7]: 1 data.shape

Out[7]: (999, 12)

In [8]: 1 data.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 999 entries, 0 to 998
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
0   user ID              999            int64
1   Name                 999            object
2   Age                  999            int64
3   Gender               999            object
4   Academic Background  999            object
5   Field of Study       999            object
6   Skills                999            object
7   Industry_Interest    999            object
8   Job_Type_Interest    999            object
9   Location_Interest    999            object
10  Salary_Expectation_(in_USD) 999            float64
11  d                    999            object

```

```
In [8]: 1 data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 999 entries, 0 to 998
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   user ID                999 non-null    int64
1   Name                   996 non-null    object
2   Age                    999 non-null    int64
3   Gender                 927 non-null    object
4   Academic Background    999 non-null    object
5   Field of Study         963 non-null    object
6   Skills                 999 non-null    object
7   Industry_Interest      999 non-null    object
8   Job_Type_Interest      999 non-null    object
9   Location_Interest      995 non-null    object
10  Salary_Expectation_(in_USD) 999 non-null    int64
11  desired_company        999 non-null    object
dtypes: int64(3), object(9)
memory usage: 93.8+ KB
```

DESCRIPTIVE STATISTICS

```
In [9]: 1 # Get descriptive statistics
2 descriptive_stats = data.describe()
3
4 # Display the descriptive statistics
5 print(descriptive_stats)
```

	user ID	Age	Salary_Expectation_(in_USD)
count	999.000000	999.000000	999.000000
mean	931435.157157	30.625626	65385.385385
std	158544.732526	9.187644	28965.576740
min	382.000000	0.000000	20000.000000
25%	962133.000000	24.000000	45000.000000
50%	977695.000000	30.000000	65000.000000

CHECKING FOR MISSING VALUES

```
In [10]: 1 # Check for missing values in the dataset
          2 missing_values = data.isnull().sum()
          3
          4 # Display the count of missing values for each column
          5 print(missing_values)
```

```
user ID          0
Name             3
Age              0
Gender           72
Academic Background  0
Field of Study   36
Skills           0
Industry_Interest  0
Job_Type_Interest  0
Location_Interest  4
Salary_Expectation_(in_USD)  0
desired_company  0
dtype: int64
```

```
In [11]: 1 # Impute missing values for numerical columns with the mean
          2 numerical_columns = data.select_dtypes(include='number').columns
          3 data[numerical_columns] = data[numerical_columns].fillna(data[numerical_columns].mean())
          4
          5 # Impute missing values for categorical columns with the most frequent value
          6 categorical_columns = data.select_dtypes(include='object').columns
          7 data[categorical_columns] = data[categorical_columns].fillna(data[categorical_columns].mode().iloc[0])
```


```
In [12]: 1 # Check for missing values in the dataset
          2 missing_values = data.isnull().sum()
          3
          4 # Display the count of missing values for each column
          5 print(missing_values)
```

```
In [12]: 1 # Check for missing values in the dataset
2 missing_values = data.isnull().sum()
3
4 # Display the count of missing values for each column
5 print(missing_values)
```

```
user ID          0
Name             0
Age              0
Gender           0
Academic Background  0
Field of Study   0
Skills           0
Industry_Interest  0
Job_Type_Interest  0
Location_Interest  0
Salary_Expectation_(in_USD)  0
desired_company  0
dtype: int64
```

```
In [13]: 1 # Plot histogram of Age
2 plt.figure(figsize=(8, 6))
3 sns.histplot(data['Age'], bins=20, kde=True)
4 plt.xlabel('Age')
5 plt.ylabel('Count')
6 plt.title('Histogram of Age')
7 plt.show()
```

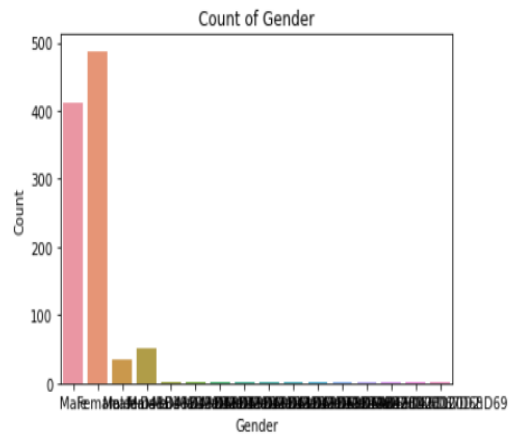
```
In [13]: 1 # Plot histogram of Age
2 plt.figure(figsize=(8, 6))
3 sns.histplot(data['Age'], bins=20, kde=True)
4 plt.xlabel('Age')
5 plt.ylabel('Count')
6 plt.title('Histogram of Age')
7 plt.show()
```



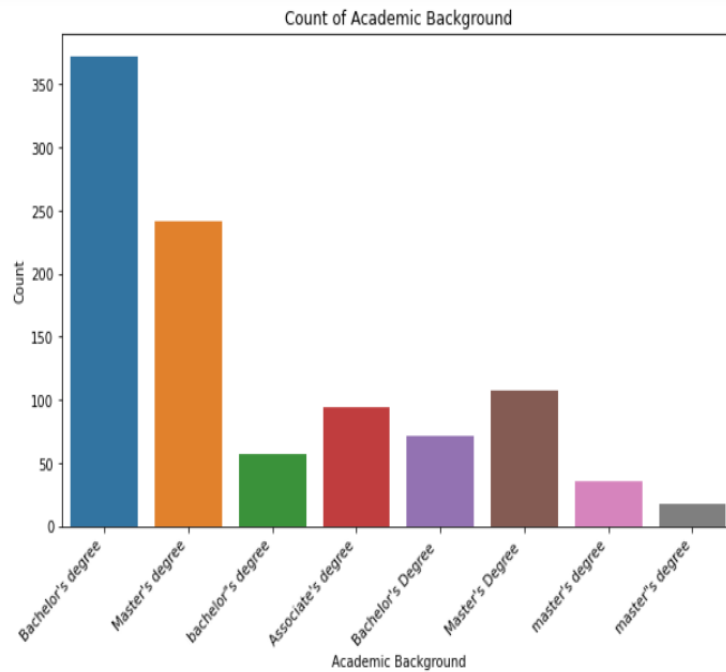
```
In [14]: 1 # Plot count of Gender
2 plt.figure(figsize=(6, 4))
3 sns.countplot(data['Gender'])
4 plt.xlabel('Gender')
5 plt.ylabel('Count')
6 plt.title('Count of Gender')
7 plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword argument: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



```
In [15]: 1 # Plot count of Academic Background
2 plt.figure(figsize=(10, 6))
3 sns.countplot(data['Academic Background'])
4 plt.xlabel('Academic Background')
5 plt.ylabel('Count')
6 plt.title('Count of Academic Background')
7 plt.xticks(rotation=45, ha='right')
8 plt.show()
```

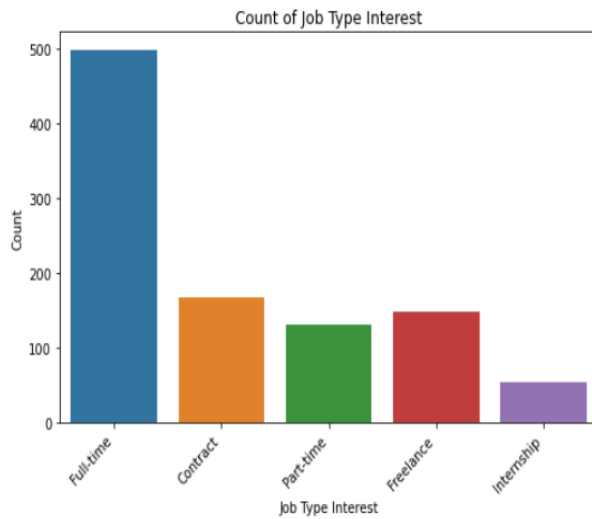



```
In [17]: 1 # Plot count of Job Type Interest
2 plt.figure(figsize=(8, 5))
3 sns.countplot(data['Job_Type_Interest'])
4 plt.xlabel('Job Type Interest')
5 plt.ylabel('Count')
6 plt.title('Count of Job Type Interest')
7 plt.xticks(rotation=45, ha='right')
8 plt.show()
```

```
In [17]: 1 # Plot count of Job Type Interest
2 plt.figure(figsize=(8, 5))
3 sns.countplot(data['Job_Type_Interest'])
4 plt.xlabel('Job Type Interest')
5 plt.ylabel('Count')
6 plt.title('Count of Job Type Interest')
7 plt.xticks(rotation=45, ha='right')
8 plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword argument: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



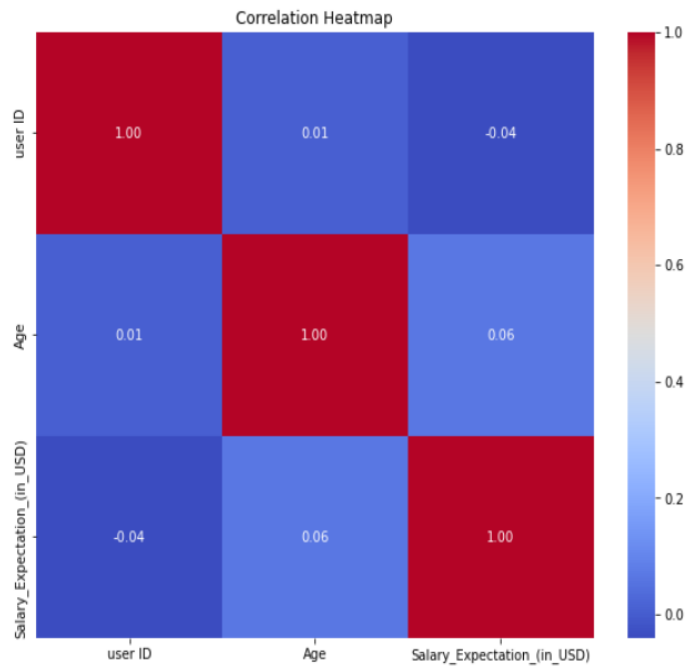
```
In [19]: 1 # Plot box plot of Salary Expectation
2 plt.figure(figsize=(8, 6))
3 sns.boxplot(data['Salary_Expectation_(in_USD)'])
4 plt.xlabel('Salary Expectation (in USD)')
5 plt.title('Box Plot of Salary Expectation')
6 plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword argument: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



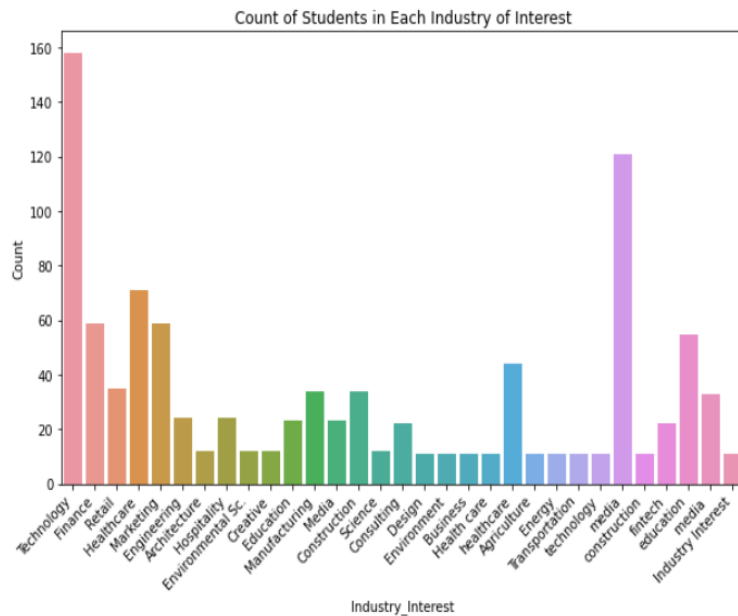
```
In [20]: 1 # Plot correlation heatmap for numerical columns
2 plt.figure(figsize=(10, 8))
3 sns.heatmap(data.corr(), annot=True, cmap='coolwarm', fmt='.2f')
4 plt.title('Correlation Heatmap')
5 plt.show()
```



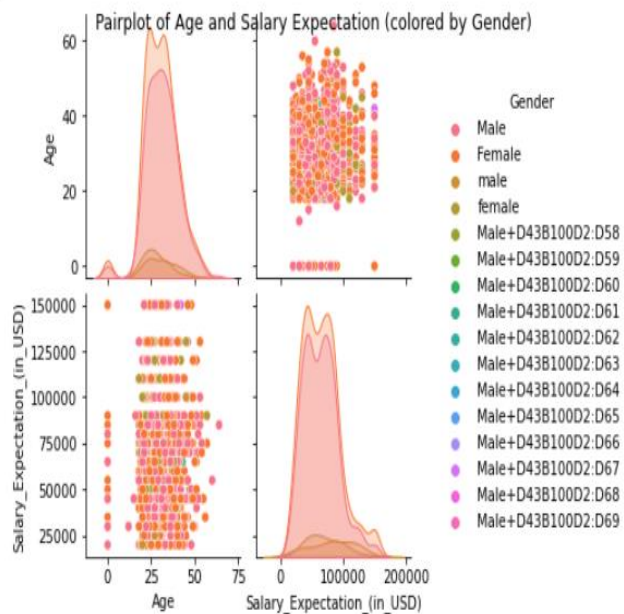
```
In [22]: 1 plt.figure(figsize=(10, 6))
2 sns.countplot(data['Industry_Interest'])
3 plt.xlabel('Industry_Interest')
4 plt.ylabel('Count')
5 plt.title('Count of Students in Each Industry of Interest')
6 plt.xticks(rotation=45, ha='right')
7 plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword argument: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



```
In [24]: 1 sns.pairplot(data, vars=['Age', 'Salary_Expectation_(in_USD)'], hue='Gender')
2 plt.suptitle('Pairplot of Age and Salary Expectation (colored by Gender)')
3 plt.show()
```



```
In [27]: 1 plt.figure(figsize=(10, 6))
2 sns.violinplot(x='Job_Type_Interest', y='Salary_Expectation_(in_USD)', data=data)
3 plt.xlabel('Job Type Interest')
4 plt.ylabel('Salary_Expectation_(in_USD)')
5 plt.title('Distribution of Salary Expectation across Job Types')
6 plt.xticks(rotation=45, ha='right')
7 plt.show()
```



```
In [30]: 1 plt.figure(figsize=(10, 6))
2 sns.boxplot(x='Academic Background', y='Salary_Expectation_(in_USD)', data=data)
3 plt.xlabel('Academic Background')
4 plt.ylabel('Salary_Expectation_(in_USD)')
5 plt.title('Comparison of Salary Expectation for Different Academic Backgrounds')
6 plt.xticks(rotation=45, ha='right')
7 plt.show()
```



```
In [32]: 1 plt.figure(figsize=(8, 6))
2 sns.scatterplot(x='Age', y='Salary_Expectation_(in_USD)', data=data, hue='Gender')
3 plt.xlabel('Age')
4 plt.ylabel('Salary_Expectation_(in_USD)')
5 plt.title('Scatter Plot of Age vs. Salary Expectation (colored by Gender)')
6 plt.show()
```




```
In [34]: 1 #One-Hot Encoding
2 one_hot_encoded_data = pd.get_dummies(data, columns=['Academic Background', 'Field of Study', 'Industry_Interest', 'Job_Type'])
3
4
5 label_encoder = LabelEncoder()
6 data['Gender_encoded'] = label_encoder.fit_transform(data['Gender'])
7
8 # Display the encoded datasets
9 print("One-Hot Encoded Data:")
10 print(one_hot_encoded_data.head())
11
12 print("\nLabel Encoded Data:")
13 print(data[['Gender', 'Gender_encoded']].head())
```

One-Hot Encoded Data:

	user ID	Name	Age	Gender \
0	986206	John Smith	28	Male
1	769632	Jane Doe	42	Female
2	981314	David Lee	35	Male
3	962892	Sarah Johnson	27	Female
4	967782	Michael Williams	46	Male

	Skills	Salary_Expectation_(in_USD) \
0	Java, Python, Data Structures	65000
1	Finance, Accounting, Microsoft Excel	85000
2	Sales, Customer Service, Communication	30000
3	Patient Care, Medical Terminology	45000
4	Digital Marketing, Social Media	70000

	Academic Background_Associate's degree \
0	0
1	0
2	0
3	1
4	0

Academic Background_Bachelor's Degree \

```

In [36]: 1 #Extracting First Name from Name
2 data['First Name'] = data['Name'].str.split().str[0]
3
4 # Feature 2: Converting 'Age' to Age Group
5 def get_age_group(age):
6     if age < 25:
7         return 'Young'
8     elif age >= 25 and age < 40:
9         return 'Middle-aged'
10    else:
11        return 'Senior'
12
13 data['Age Group'] = data['Age'].apply(get_age_group)
14
15 # Feature 3: Counting the number of skills each student has
16 data['Number of Skills'] = data['Skills'].apply(lambda x: len(x.split(',')))
17
18 # Feature 4: Encoding 'Full-time' as 1 and 'Part-time' as 0 for Job Type Interest
19 data['Job_Type_Interest'] = data['Job_Type_Interest'].apply(lambda x: 1 if x == 'Full-time' else 0)
20
21 # Feature 5: Encoding 'Male' as 1 and 'Female' as 0 for Gender
22 data['Gender'] = data['Gender'].apply(lambda x: 1 if x == 'Male' else 0)
23
24 # Display the updated dataset with engineered features
25 print(data.head())

```

```

      Field of Study      Skills \
0  Computer Science 2      Java, Python, Data Structures
1           Business      Finance, Accounting, Microsoft Excel
2  Computer Science      Sales, Customer Service, Communication
3           Nursing      Patient Care, Medical Terminology
4           Marketing      Digital Marketing, Social Media

      Industry_Interest  Job_Type_Interest  Location_Interest \
0           Technology              1      New York
1           Finance              0      London
2           Retail              0      Chicago
3           Healthcare              1      Los Angeles
4           Marketing              0      Toronto

```

```

In [37]: 1 # Prepare the feature matrix X and target vector y
2 X = data[['Age', 'Gender']] # Select relevant features
3 y = data['Job_Type_Interest'] # Target variable
4
5 # Split the data into training and testing sets (80% train, 20% test)
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
7
8 # Initialize and train the KNN classifier
9 knn_classifier = KNeighborsClassifier(n_neighbors=5)
10 knn_classifier.fit(X_train, y_train)
11
12 # Make predictions on the test set
13 y_pred = knn_classifier.predict(X_test)
14
15 # Evaluate the classifier's performance
16 accuracy = accuracy_score(y_test, y_pred)
17 print("KNN Classifier Accuracy:", accuracy)

```

KNN Classifier Accuracy: 0.525

```

In [45]: 1 # Import required libraries
2 from surprise import Dataset, Reader
3 from surprise.model_selection import train_test_split
4 from surprise import KNNBasic
5 from surprise.accuracy import rmse
6
7 # Prepare the data for Surprise library
8 reader = Reader(rating_scale=(0, 1))
9 surprise_data = Dataset.load_from_df(data[['user ID', 'Salary_Expectation_(in_USD)', 'Age']], reader)
10
11 # Split the data into training and testing sets (80% train, 20% test)
12 trainset, testset = train_test_split(surprise_data, test_size=0.2, random_state=42)
13
14 # Initialize and train the KNNBasic collaborative filtering model
15 knn_collaborative = KNNBasic(sim_options={'user_based': True})
16 knn_collaborative.fit(trainset)
17
18 # Make predictions on the test set
19 predictions = knn_collaborative.test(testset)

```

```

In [45]: 1 # Import required libraries
2 from surprise import Dataset, Reader
3 from surprise.model_selection import train_test_split
4 from surprise import KNNBasic
5 from surprise.accuracy import rmse
6
7 # Prepare the data for Surprise library
8 reader = Reader(rating_scale=(0, 1))
9 surprise_data = Dataset.load_from_df(data[['user ID', 'Salary_Expectation_(in_USD)', 'Age']], reader)
10
11 # Split the data into training and testing sets (80% train, 20% test)
12 trainset, testset = train_test_split(surprise_data, test_size=0.2, random_state=42)
13
14 # Initialize and train the KNNBasic collaborative filtering model
15 knn_collaborative = KNNBasic(sim_options={'user_based': True})
16 knn_collaborative.fit(trainset)
17
18 # Make predictions on the test set
19 predictions = knn_collaborative.test(testset)
20
21 # Evaluate the model's performance (Root Mean Squared Error, RMSE)
22 rmse_score = rmse(predictions)
23 print("Collaborative Filtering RMSE:", rmse_score)

```

Computing the msd similarity matrix...

Done computing similarity matrix.

RMSE: 31.6624

Collaborative Filtering RMSE: 31.662438314191785

```

In [49]: 1 # Import required libraries
2 from surprise import Dataset, Reader
3 from surprise.model_selection import train_test_split
4 from surprise import KNNBasic
5 from surprise.accuracy import rmse
6
7 # Prepare the data for Surprise library
8 reader = Reader(rating_scale=(0, 1))
9 surprise_data = Dataset.load_from_df(data[['user ID', 'Salary_Expectation_(in_USD)', 'Age']], reader)
10
11 # Split the data into training and testing sets (80% train, 20% test)

```

```

11 # Split the data into training and testing sets (80% train, 20% test)
12 trainset, testset = train_test_split(surprise_data, test_size=0.2, random_state=42)
13
14 # Initialize and train the KNNBasic collaborative filtering model
15 knn_collaborative = KNNBasic(sim_options={'user_based': True})
16 knn_collaborative.fit(trainset)
17
18 # Make predictions on the test set
19 predictions = knn_collaborative.test(testset)
20
21 # Evaluate the model's performance (Root Mean Squared Error, RMSE)
22 rmse_score = rmse(predictions)
23 print("Collaborative Filtering RMSE:", rmse_score)

```

Computing the msd similarity matrix...

Done computing similarity matrix.

RMSE: 31.6624

Collaborative Filtering RMSE: 31.662438314191785

```

In [51]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.neighbors import KNeighborsClassifier
4 from sklearn.ensemble import RandomForestClassifier
5 from sklearn.metrics import roc_curve, auc
6
7 # Prepare the feature matrix X and target vector y
8 X = data[['Age', 'Gender', 'Number of Skills']] # Select relevant features
9 y = data['Job_Type_Interest'] # Target variable
10
11 # Split the data into training and testing sets (80% train, 20% test)
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
13
14 # Standardize the feature data for better training performance (not necessary for some algorithms)
15 scaler = StandardScaler()
16 X_train_scaled = scaler.fit_transform(X_train)
17 X_test_scaled = scaler.transform(X_test)
18
19 # Initialize models for different algorithms
20 knn_classifier = KNeighborsClassifier(n_neighbors=5)
21 knn_classifier.fit(X_train_scaled, y_train)

```

```

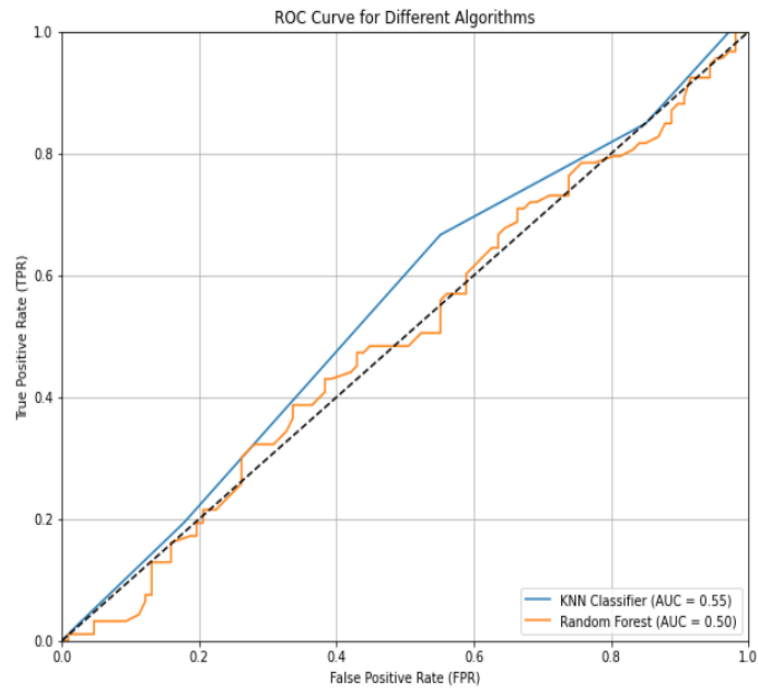
19 # Initialize models for different algorithms
20 knn_classifier = KNeighborsClassifier(n_neighbors=5)
21 knn_classifier.fit(X_train_scaled, y_train)
22
23 # Random Forest
24 rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
25 rf_classifier.fit(X_train_scaled, y_train)
26
27 # Calculate ROC curves and AUC for each model
28 models = [knn_classifier, rf_classifier]
29 model_names = ['KNN Classifier', 'Random Forest']
30
31 plt.figure(figsize=(10, 8))
32
33 for model, name in zip(models, model_names):
34     y_pred_prob = model.predict_proba(X_test_scaled)[: , 1]
35     fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
36     auc_score = auc(fpr, tpr)
37     plt.plot(fpr, tpr, label=f'{name} (AUC = {auc_score:.2f})')
38
39 plt.plot([0, 1], [0, 1], 'k--')
40 plt.xlim([0.0, 1.0])
41 plt.ylim([0.0, 1.0])
42 plt.xlabel('False Positive Rate (FPR)')
43 plt.ylabel('True Positive Rate (TPR)')
44 plt.title('ROC Curve for Different Algorithms')
45 plt.legend(loc='lower right')
46 plt.grid()
47 plt.show()

```

```

38 plt.plot([0, 1], [0, 1], 'k--')
39 plt.xlim([0.0, 1.0])
40 plt.ylim([0.0, 1.0])
41 plt.xlabel('False Positive Rate (FPR)')
42 plt.ylabel('True Positive Rate (TPR)')
43 plt.title('ROC Curve for Different Algorithms')
44 plt.legend(loc='lower right')
45 plt.grid()
46 plt.show()

```



```
In [52]: 1 accuracy = accuracy_score(y_test, y_pred_prob.round()) # Calculate accuracy using y_pred_prob
2 report = classification_report(y_test, knn_classifier.predict(X_test_scaled)) # Using the KNN classifier
3
4 print(f"--- {name} ---")
5 print("Accuracy:", accuracy)
6 print("Classification Report:")
7 print(report)
8 print("-----")
```

--- Random Forest ---

Accuracy: 0.485

Classification Report:

	precision	recall	f1-score	support
0	0.61	0.45	0.52	107
1	0.51	0.67	0.58	93
accuracy			0.55	200
macro avg	0.56	0.56	0.55	200
weighted avg	0.56	0.55	0.55	200

```
In [53]: 1 # Get feature importances from the trained Random Forest classifier
2 feature_importances = rf_classifier.feature_importances_
3
4 # Print feature importances
5 print("Feature Importances:")
6 for feature, importance in zip(X.columns, feature_importances):
7     print(f"{feature}: {importance}")
```

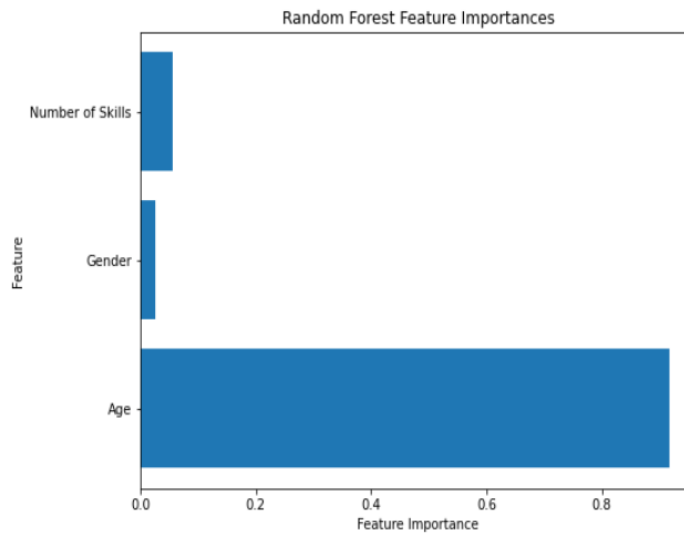
Feature Importances:

Age: 0.9170870182907975

Gender: 0.025868912623061253

Number of Skills: 0.05704406908614126


```
In [54]: 1 # Get feature importances from the trained Random Forest classifier
2 feature_importances = rf_classifier.feature_importances_
3
4 # Plot feature importances in a bar plot
5 feature_names = X.columns
6 plt.figure(figsize=(8, 6))
7 plt.barh(feature_names, feature_importances)
8 plt.xlabel('Feature Importance')
9 plt.ylabel('Feature')
10 plt.title('Random Forest Feature Importances')
11 plt.show()
```



```
In [57]: 1 import joblib
2 import os
3 from sklearn.preprocessing import OneHotEncoder, StandardScaler
4
5 # Create the directory
6 output_directory = 'student_job_recommendation_system/'
7 os.makedirs(output_directory, exist_ok=True)
8
9 # Assuming you have defined the 'academic_data' variable before this point
10 # Convert categorical variables to numerical using one-hot encoding
11 encoder = OneHotEncoder()
12 academic_encoded = encoder.fit_transform(data).toarray()
13
14 # Scale the data for better clustering performance
15 scaler = StandardScaler()
16 academic_scaled = scaler.fit_transform(academic_encoded)
17
18 # Save the trained machine learning models
19 # Assuming you already have the trained models: knn_classifier, knn_collaborative, bayesian_model, model, rf_classifier
20 joblib.dump(knn_classifier, 'student_job_recommendation_system/knn_classifier_model.pkl')
21 joblib.dump(rf_classifier, 'student_job_recommendation_system/random_forest_model.pkl')
```

Out[57]: ['student_job_recommendation_system/random_forest_model.pkl']